

Московский государственный технический университет им. Н.Э. Баумана  
Факультет «Информатика и системы управления»  
Кафедра «Системы обработки информации и управления»



**Отчёт**

**“Методы машинного обучения”**

**Лабораторная работа № 2**

**“Изучение библиотек обработки данных”**

ИСПОЛНИТЕЛЬ:

Студент группы ИУ5-21М

Коростелёв В. М. \_\_\_\_\_

ПРЕПОДАВАТЕЛЬ:

Гапанюк Ю. Е. \_\_\_\_\_

**Москва – 2019**

---

# Задание

Часть 1. Выполните первое демонстрационное задание "demo assignment" под названием "Exploratory data analysis with Pandas" со страницы курса <https://mlcourse.ai/assignments> Часть 2. Выполните следующие запросы с использованием двух различных библиотек - Pandas и PandaSQL: один произвольный запрос на соединение двух наборов данных один произвольный запрос на группировку набора данных с использованием функций агрегирования Сравните время выполнения каждого запроса в Pandas и PandaSQL.

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: data = pd.read_csv('adult.data.csv')
data.head()
```

Out[2]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	21790
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0

```
In [3]: data['sex'].value_counts()
```

Out[3]: Male 21790
Female 10771
Name: sex, dtype: int64

```
In [4]: data.loc[data['sex'] == 'Female', 'age'].mean()
```

Out[4]: 36.85823043357163

```
In [5]: print("{0:%}".format(data[data["native-country"] == "Germany" ].shape[0] / data.shape[0]))
```

0.420749%

```
In [6]: ages1 = data[data["salary"] == "<=50K"]["age"]
ages2 = data[data["salary"] == ">50K"]["age"]
print("<=50K: = {0} ± {1} years".format(ages1.mean(), ages1.std()))
print(">50K: = {0} ± {1} years".format(ages2.mean(), ages2.std()))
```

<=50K: = 36.78373786407767 ± 14.02008849082488 years  
>50K: = 44.24984058155847 ± 10.519027719851826 years

```
In [7]: high_educations = set(["Bachelors", "Prof-school", "Assoc-acdm", "Assoc-voc", "Masters", "Doctorate"])
def high_educated(e):
    return e in high_educations
data[data["salary"] == ">50K"]["education"].map(high_educated).all()
```

Out[7]: False

```
In [8]: data.groupby(["race", "sex"])["age"].describe()
```

Out[8]:

		count	mean	std	min	25%	50%	75%	max
race	sex								
Amer-Indian-Eskimo	Female	119.0	37.117647	13.114991	17.0	27.0	36.0	46.00	80.0
	Male	192.0	37.208333	12.049563	17.0	28.0	35.0	45.00	82.0
Asian-Pac-Islander	Female	346.0	35.089595	12.300845	17.0	25.0	33.0	43.75	75.0
	Male	693.0	39.073593	12.883944	18.0	29.0	37.0	46.00	90.0
Black	Female	1555.0	37.854019	12.637197	17.0	28.0	37.0	46.00	90.0
	Male	1569.0	37.682600	12.882612	17.0	27.0	36.0	46.00	90.0
Other	Female	109.0	31.678899	11.631599	17.0	23.0	29.0	39.00	74.0
	Male	162.0	34.654321	11.355531	17.0	26.0	32.0	42.00	77.0
White	Female	8642.0	36.811618	14.329093	17.0	25.0	35.0	46.00	90.0
	Male	19174.0	39.652498	13.436029	17.0	29.0	38.0	49.00	90.0

```
In [9]: data[(data["race"] == "Amer-Indian-Eskimo")
            & (data["sex"] == "Male")]["age"].max()
```

Out[9]: 82

```
In [10]: def is_married(m):
          return m.startswith("Married")
data["married"] = data["marital-status"].map(is_married)
(data[(data["sex"] == "Male") & (data["salary"] == ">50K")]
 ["married"]).value_counts()
```

Out[10]: True 5965  
False 697  
Name: married, dtype: int64

```
In [11]: m = data["hours-per-week"].max()
print("Maximum is {} hours/week.".format(m))
people = data[data["hours-per-week"] == m]
c = people.shape[0]
print("{} people work this time at week.".format(c))
s = people[people["salary"] == ">50K"].shape[0]
print("{0:%} get >50K salary.".format(s / c))
```

Maximum is 99 hours/week.  
85 people work this time at week.  
29.411765% get >50K salary.

```
In [12]: p = pd.crosstab(data["native-country"], data["salary"],
                        values=data['hours-per-week'], aggfunc="mean")
p
```

Out[12]:

	salary	<=50K	>50K
native-country			
	?	40.164760	45.547945
	Cambodia	41.416667	40.000000
	Canada	37.914634	45.641026
	China	37.381818	38.900000
	Columbia	38.684211	50.000000
	Cuba	37.985714	42.440000
	Dominican-Republic	42.338235	47.000000
	Ecuador	38.041667	48.750000
	El-Salvador	36.030928	45.000000
	England	40.483333	44.533333
	France	41.058824	50.750000
	Germany	39.139785	44.977273
	Greece	41.809524	50.625000
	Guatemala	39.360656	36.666667
	Haiti	36.325000	42.750000
	Holand-Netherlands	40.000000	NaN
	Honduras	34.333333	60.000000
	Hong	39.142857	45.000000
	Hungary	31.300000	50.000000
	India	38.233333	46.475000
	Iran	41.440000	47.500000
	Ireland	40.947368	48.000000
	Italy	39.625000	45.400000
	Jamaica	38.239437	41.100000
	Japan	41.000000	47.958333
	Laos	40.375000	40.000000
	Mexico	40.003279	46.575758
	Nicaragua	36.093750	37.500000
	Outlying-US(Guam-USVI-etc)	41.857143	NaN
	Peru	35.068966	40.000000
	Philippines	38.065693	43.032787
	Poland	38.166667	39.000000
	Portugal	41.939394	41.500000
	Puerto-Rico	38.470588	39.416667
	Scotland	39.444444	46.666667
	South	40.156250	51.437500
	Taiwan	33.774194	46.800000
	Thailand	42.866667	58.333333
	Trinidad&Tobago	37.058824	40.000000
	United-States	38.799127	45.505369

	salary	<=50K	>50K
native-country			
	Vietnam	37.193548	39.200000
	Yugoslavia	41.600000	49.500000

In [13]: `p.loc["Japan"]`

Out[13]: salary  
<=50K 41.000000  
>50K 47.958333  
Name: Japan, dtype: float64

In [14]: `from pandasql import sqldf`  
`pysqldf = lambda q: sqldf(q, globals())`

In [15]: `wind = (pd.read_csv('wind speed.csv', header=None, names=["row", "UNIX", "date", "time", "speed", "text"]).drop("text", axis=1))`  
`temp = (pd.read_csv('temperature.csv', header=None, names=["row", "UNIX", "date", "time", "temperature", "text"]).drop("text", axis=1))`

In [16]: `wind.head()`

Out[16]:

	row	UNIX	date	time	speed
0	1	1475315718	2016-09-30	23:55:18	7.87
1	2	1475315423	2016-09-30	23:50:23	7.87
2	3	1475315124	2016-09-30	23:45:24	9.00
3	4	1475314821	2016-09-30	23:40:21	13.50
4	5	1475314522	2016-09-30	23:35:22	15.75

In [17]: `wind.dtypes`

Out[17]: row int64  
UNIX int64  
date object  
time object  
speed float64  
dtype: object

In [18]: `temp.head()`

Out[18]:

	row	UNIX	date	time	temperature
0	1	1475315718	2016-09-30	23:55:18	48
1	2	1475315423	2016-09-30	23:50:23	48
2	3	1475315124	2016-09-30	23:45:24	48
3	4	1475314821	2016-09-30	23:40:21	48
4	5	1475314522	2016-09-30	23:35:22	48

```
In [19]: temp.dtypes
```

```
Out[19]: row          int64
UNIX          int64
date          object
time          object
temperature   int64
dtype: object
```

```
In [20]: wind.merge(temp[["UNIX", "temperature"]], on="UNIX").head()
```

```
Out[20]:
```

	row	UNIX	date	time	speed	temperature
0	1	1475315718	2016-09-30	23:55:18	7.87	48
1	2	1475315423	2016-09-30	23:50:23	7.87	48
2	3	1475315124	2016-09-30	23:45:24	9.00	48
3	4	1475314821	2016-09-30	23:40:21	13.50	48
4	5	1475314522	2016-09-30	23:35:22	15.75	48

```
In [21]: %%timeit
wind.merge(temp[["UNIX", "temperature"]], on="UNIX")
```

20.6 ms ± 2.41 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

```
In [22]: pysqldf("""SELECT w.row, w.UNIX, w.date, w.time,
w.speed, t.temperature
FROM wind AS w JOIN temp AS t
ON w.UNIX = t.UNIX """).head()
```

```
Out[22]:
```

	row	UNIX	date	time	speed	temperature
0	1	1475315718	2016-09-30	23:55:18	7.87	48
1	2	1475315423	2016-09-30	23:50:23	7.87	48
2	3	1475315124	2016-09-30	23:45:24	9.00	48
3	4	1475314821	2016-09-30	23:40:21	13.50	48
4	5	1475314522	2016-09-30	23:35:22	15.75	48

```
In [23]: %%timeit
pysqldf("""SELECT w.row, w.UNIX, w.date, w.time,
w.speed, t.temperature
FROM wind AS w JOIN temp AS t
ON w.UNIX = t.UNIX """).head()
```

758 ms ± 17.6 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
In [24]: wind.groupby("date")["speed"].mean().head()
```

```
Out[24]: date
2016-09-01    6.396560
2016-09-02    5.804086
2016-09-03    4.960248
2016-09-04    5.184571
2016-09-05    5.830676
Name: speed, dtype: float64
```

```
In [25]: %%timeit
wind.groupby("date")["speed"].mean().head()
```

3.36 ms ± 75.6 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)

```
In [26]: pysqldf("""SELECT date, AVG(speed) FROM wind GROUP BY date """).head()
```

Out[26]:

	date	AVG(speed)
0	2016-09-01	6.396560
1	2016-09-02	5.804086
2	2016-09-03	4.960248
3	2016-09-04	5.184571
4	2016-09-05	5.830676

```
In [27]: %%timeit
pysqldf("""SELECT date, AVG(speed) FROM wind GROUP BY date """).head()
```

293 ms ± 8.32 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)